

Loom: Funded Paired Trial

Decisions taken at the close of 2026-09-18, marked to market on 2026-09-25. 1000 USD notional per arm, six arms, 42 positions.

Result

Loom reduced returns in all three cohorts. No arm improved. This is the first time Loom has been scored on money rather than on whether it changed a decision, and it lost.

Arm	Deployed	Cash	Final	Return	vs QQQ
gambler-control	1000	0	1049.28	+4.93%	+1.63%
gambler-loom	1000	0	1036.49	+3.65%	+0.35%
professional-control	1000	0	1030.85	+3.09%	-0.22%
amateur-control	900	100	1017.95	+1.80%	-1.51%
professional-loom	950	50	1011.55	+1.16%	-2.15%
amateur-loom	650	350	999.86	-0.01%	-3.32%
QQQ benchmark			1033.02	+3.30%	

Paired, with Loom as the only variable

Cohort	Control	With Loom	Difference
Professional	+3.09%	+1.16%	-1.93 pts
Amateur	+1.80%	-0.01%	-1.81 pts
Gambler	+4.93%	+3.65%	-1.28 pts

Why, mechanically

The week rose +3.30%. Loom's consistent and reproducible effect across all three trials run so far is to reduce risk taking: less capital deployed, more passes, smaller positions. In a rising week that is a cost by construction, and this quantifies it at roughly one and a third to two points.

Company	Move over the week
AAPL	+1.47%
AMD	+12.65%
AMZN	-1.59%
GOOGL	-1.61%
MSFT	+4.53%
SNDK	-0.78%
TSLA	+2.15%

The decisive position was AMD at +12.6%, the best performer of the seven. Both the professional and the amateur Loom arms passed on it, citing insider selling that Loom had surfaced. That single call accounts for most of the shortfall. It was a real and checkable signal and over five sessions it was wrong.

The amateur arm with Loom held 350 of its 1000 in cash. Loom talked it out of deploying, and cash earns nothing in an up week.

What this shows, and what it does not

It shows that Loom's caution has a price, and that the price is not small. Every previous trial recorded 'capital deployed fell' as a neutral observation. It is not neutral. Measured over this week it costs between 1.28 and 1.93 points.

It does not show that Loom destroys value. One week, five sessions, seven companies and three pairs is overwhelmingly noise. The sign would be expected to invert in a falling week, which is the entire purpose of reducing exposure. This test cannot distinguish a tool that trades upside for downside protection from one that simply costs one and a half points, and claiming otherwise from three paired observations would be the kind of overreach the rest of this project exists to avoid.

Method

Decisions were taken as at the close on 2026-09-18 and every input was rebuilt to that date. Factor rankings were recomputed with the point in time scorer rather than read from today's stored table; findings were filtered to those dated on or before the decision date; price context and market reaction windows were closed before it. The packets were audited for any date after the decision date and none was found. Verified independently: Microsoft reads as down two percent over the year in the packet and is up two percent today, so the rebuild took.

Each arm received a notional 1000 units and allocated across long, short or no position, with leftover cash a permitted and meaningful choice. Positions were marked at the adjusted close on 2026-09-25. The benchmark is QQQ over the identical window.

Flaws in this run, disclosed

The agents had read access to a live API. They needed it to record decisions, and the same access serves current prices. One arm quoted end of week closes in its report. Every stored rationale was searched for those figures and none contained them, and that arm's book placed its second largest position in the week's second worst performer and its smallest in the second best, which is not what foreknowledge produces. The hole was nonetheless real and should be closed by having agents write to a file rather than to the API.

Two agents initially refused the task because the prompt described a 1000 USD budget as real money. It is a paper simulation and the wording was simply false. They were right to decline and were re-run with accurate framing.

The horizon is too short to conclude anything. Five sessions is noise. A three month window is the minimum that would begin to separate signal from weather.

Limitations

Three paired observations. Agents are language models rather than people. The seven companies are among the best covered in the database and are not representative of the remaining nine hundred and ninety nine. No significance is claimed for any figure here, and none should be read into the ordering of the arms.