

Loom: Funded Paired Trial, One Quarter

Books set at the close of 2026-06-25, marked on 2026-09-25. 65 trading sessions. 1000 units notional per arm, six arms, 42 positions.

Result

The result splits, and the split is the finding. Loom reduced returns for the two arms that were already doing well and rescued the one that was losing money. It turned the only negative book into a positive one, an 8.52 point swing and the largest single effect measured in any trial of Loom so far.

Arm	Deployed	Cash	Final	Return	vs QQQ
professional-control	800	200	1173.82	+17.38%	+13.35%
professional-loom	850	150	1136.34	+13.63%	+9.60%
amateur-control	850	150	1102.97	+10.30%	+6.26%
amateur-loom	650	350	1095.14	+9.51%	+5.48%
gambler-loom	950	50	1045.18	+4.52%	+0.48%
gambler-control	850	150	959.97	-4.00%	-8.04%
QQQ benchmark			1040.34	+4.03%	

Paired, with Loom as the only variable

Cohort	Control	With Loom	Difference	
Professional	+17.38%	+13.63%	-3.75 pts	hurt
Amateur	+10.30%	+9.51%	-0.78 pts	hurt
Gambler	-4.00%	+4.52%	+8.52 pts	helped

The quarter

Company	Move
MSFT	+46.57%
AAPL	+24.06%
AMD	+18.41%
AMZN	+9.98%
GOOGL	+0.13%
TSLA	-0.80%
SNDK	-23.86%

A seventy point spread between best and worst, against a benchmark of +4.03%. Stock selection could express itself, which is why this window was chosen over the one week run that preceded it.

Why the split

The gambler is the case Loom is built for

Its control committed 40 percent to SNDK on momentum alone: up 4842 percent over the year and trading near its high. SNDK then fell 23.86 percent. The arm with Loom read the same names and halved that position to 20 percent, citing the price against value contradiction and the valuation percentiles. In its own words, the warning that the share had run hard and was now expensive is what stopped it oversizing. It did not obey Loom; it was tempered by it, and that half sized position is most of the 8.52 points.

The professional lost by being hedged, not by being wrong

Its Loom book ran 31 percent net long against the control's 80, using SNDK and AMD shorts as hedges. The SNDK short worked and the AMD short did not, AMD returning 18.41 percent, and the reduced net exposure gave up most of a rising quarter. Loom's information was largely correct and the risk posture it encouraged was expensive. Worth recording: that arm upsized Microsoft to 20 percent on a Loom finding about committed future orders that did not appear in the control packet at all, and Microsoft returned 46.57 percent. Loom's single best call of the quarter, and the book still trailed.

The amateur is close to a wash

At 0.78 points behind, the difference is smaller than the noise. The behavioural signature is the same as the one week run: it held 350 of its 1000 in cash against the control's 150. Loom makes beginners cautious, and caution costs in a rising quarter.

What now holds across both funded runs

Loom reduces variance rather than raising return. It hurt the two arms that were already doing well and rescued the one that was blowing up. Over one week, when nothing was blowing up, it only cost. Over a quarter containing a genuine bubble, it paid for itself once in three.

That is a coherent story and it is not the story the product tells. Nothing in Loom's interface says it will cost you in a rally and save you in a blow up. If this pattern survives more observations, it is the single most important thing Loom should say about itself, and it changes who the product is for: not someone trying to beat the market, but someone trying not to lose badly.

Method, and the fixes applied since the one week run

Every input was rebuilt to the close on 2026-06-25. Factor rankings were recomputed with the point in time scorer rather than read from the stored table. Findings were filtered to those dated on or before the decision date. Price context and market reaction windows were closed before it. The packets were audited for dates after the decision date and for any of the seven closing prices from the end of the window, and both searches came back empty.

Three flaws from the one week run were corrected:

Agents had read access to a live API serving current prices. Removed entirely. Agents write a JSON file and the harness ingests it, so no endpoint can leak the future to them.

The prompt described a notional budget as real money, and two agents correctly refused. Reframed accurately as a paper simulation. No refusals.

A five session horizon is noise. Extended to 65 sessions, with a seventy point spread between the best and worst name.

One judgement is worth stating plainly. Findings dated before the decision date were used even though Loom itself read those documents later, because the documents were public at the time and the question is whether the information was available, not when this particular reader got to it. A stricter test would also hold Loom's own reading schedule fixed, and would have almost no findings to work with.

Limitations

Three paired observations in this run and six across both funded runs. No significance is claimed and none should be inferred from the ordering of the arms. The agents are language models rather than people. The seven companies are among the best covered in the database and are not representative of the remaining nine hundred and ninety nine. One quarter contains one regime, and this one happened to contain a bubble deflating, which is the regime most favourable to a tool that warns about expensive momentum. A quarter in which momentum kept working would likely reverse the sign on the gambler.

The decisions in this run are stored with their forward, benchmark and abnormal returns computed, which is the first time any Loom trial has been scored on outcome rather than on whether a decision changed.