

Loom: Paired Agent Trial

Run trial-2026-09-27 | six arms, three personas, seven companies, 42 decisions | 2026-09-27

Summary

Loom changed **8 of 21** paired decisions. Capital committed fell in every cohort. One decision reversed outright, from a long to a short. All three Loom arms independently declined Alphabet, each citing Loom's explicit refusal to form a view, which reproduces the single clearest result of the September trial with a different underlying model.

The uncomfortable finding is the opposite of September's. Then, not one agent cited Loom's directional verdict and every changed decision was changed by a fact. Now the verdict is referenced in eight of twenty-one rationales and is visibly driving decisions. It remains the one component with no demonstrated predictive validity across five separate measurement approaches, so its new prominence is a regression to watch rather than a success.

Method

Three personas, each run twice: once with only baseline market data, once with that same data plus Loom's case file. The personas are a professional equity analyst, a complete beginner, and a retail speculator. Neither arm was told the hypothesis. Loom was described to the treatment arms as one input among others, which they were free to ignore, and each was asked to record which specific items it actually used. An empty list was stated to be a valid answer.

Seven companies, unchanged from September so the universe is held constant: AAPL, AMD, AMZN, GOOGL, MSFT, SNDK, TSLA. Each agent returned an action, a fraction of capital, a conviction, a written rationale of at least forty characters, and what would change its mind. Decisions were submitted through the same endpoint as the first trial and are stored beside it.

What changed in Loom between the two trials

The first trial diagnosed five defects. All five were addressed before this run, so the rerun is partly a test of those fixes:

Stance was substantially a headcount of findings, so thorough disclosure read as bad news. Replaced with a residual against what documents of the same genre normally carry.

Thirty to sixty findings withheld per company with no way to reach them. Served in full, ranked, with routine boilerplate tagged.

One contradiction type defeated a reader outright. Plain explanations written for all six, barred from using terms of art.

Single-document reads produced uniformly negative verdicts. Informative-evidence guard; a record made only of boilerplate now refuses.

Nobody cited the verdict; every changed decision came from a fact. Verdict demoted below the basis line and relabelled to state what it measures.

Comparability

Absolute figures cannot be compared across the two trials. The agents here ran on a smaller, cheaper model, and the control arms behaved very differently: the professional control committed 2.70 units of capital against an average of 0.09 per decision in September. That is the model, not Loom. What remains valid is the paired comparison inside this run, where persona, companies and model are held constant and Loom is the only variable. That is the comparison the design exists to make.

Results

Paired decisions

Company	Professional	Amateur	Gambler
AAPL	pass 0.00 -> pass 0.00	buy 0.15 -> buy 0.12	pass 0.00 -> pass 0.00
AMD	sell 0.40 -> pass 0.00	pass 0.00 -> pass 0.00	buy 0.80 -> buy 0.60
AMZN	buy 0.60 -> buy 0.25	buy 0.20 -> pass 0.00	pass 0.00 -> pass 0.00
GOOGL	buy 0.65 -> pass 0.00	buy 0.18 -> pass 0.00	pass 0.00 -> pass 0.00
MSFT	pass 0.00 -> buy 0.25	pass 0.00 -> pass 0.00	pass 0.00 -> pass 0.00
SNDK	short 0.50 -> sell 0.15	sell 0.05 -> pass 0.00	buy 0.90 -> buy 0.75
TSLA	buy 0.55 -> short 0.15	pass 0.00 -> pass 0.00	pass 0.00 -> pass 0.00

Action and fraction of capital, control on the left of each arrow and Loom on the right. Bold marks a changed action.

Capital committed

Cohort	Control	With Loom	Change
Professional	2.70	0.80	-1.90
Amateur	0.58	0.12	-0.46
Gambler	1.70	1.35	-0.35

Summed across the seven companies. Every cohort committed less with Loom, which reproduces September's direction on all three.

What the Loom arms cited

Of twenty-one decisions taken with Loom, **1** cited nothing from it. Individual findings dominated everything else by roughly eighteen to one. Factor rankings, valuation measures and the price context were each cited, but sparingly.

Evaluation

What held up

The refusal is still Loom's most useful output. All three Loom arms declined Alphabet and all three named the refusal as the reason. The professional control had committed 0.65 to it and the amateur control 0.18. Three readers of differing sophistication, on a different model from September, reached the same conclusion from Loom saying it did not know.

Verbatim, the beginner: Loom says there is insufficient data to form an opinion about Google. There are lots of concerns about antitrust lawsuits and regulatory stuff I do not understand. I would rather invest in companies w

The routine tag worked as designed. The speculator explicitly discarded "routine boilerplate risks, supply chain, tariffs" while using the momentum and margin material. That tag did not exist in September, and the finding it suppresses is exactly the class that made the old stance a headcount.

The genre correction is visible in decisions. The professional chose Microsoft as its best risk and reward, citing the fourteen point nine percent post-filing market reaction. Microsoft was mislabelled negative by the old counting stance and is positive under the residual.

One outright reversal. The professional moved from buying Tesla at 0.55 to shorting it at 0.15.

Verbatim: Loom completely negative: 14 concerns vs 0 positives. Robotaxi regulatory risk asymmetric (only 2% of findings this severe for company). Margin pressures from financing subvention and regulatory credit loss real. Down 16% YTD, fur

What did not

The verdict became load-bearing. Eight of twenty-one rationales reference it, against zero in September. The relabelling from "Leaning positive" to "Reads better than usual" was intended to make the claim honest about what it measures. It appears also to have made it more quotable. Loom's direction has now failed to predict returns under five separate measurement approaches, so a reader leaning on it is leaning on the weakest thing in the product.

The beginner still drowns in the quantitative half. Asked what it could not follow, the amateur arm listed operating and gross margins, the distinction between reported profit and cash, return on assets, the valuation yields, and balance sheet expansion. One item is a plain defect rather than a knowledge gap: it could not tell whether a ranking of "worst" was good or bad. The polarity of a percentile is not legible without knowing the direction of the underlying measure, and Loom does not state it.

Two subsystems contradict each other on the page. On Apple, the loss of Google search payments is presented as the best argument the other way and is simultaneously tagged routine. Both are correct and they measure different things: the statistical engine finds the severity rare for Apple, the genre norms find the direction uninformative for a risk factor. On one row they read as a contradiction.

What Loom needs

- 1. Score these decisions against what actually happened.** Both trials measured whether Loom changed decisions. Neither measured whether it improved them. The schema already carries forward, benchmark and abnormal return columns and they are empty. Scoring 84 decisions across two runs in several weeks is the only test that would tell whether any of this helps, and it costs nothing but time.
- 2. Decide what the verdict is for.** It is now prominent and still unvalidated. Two coherent options: demote it much harder so the facts carry the page, or keep it and accept that Loom leads with its weakest claim. The current position, honestly labelled but increasingly relied upon, is the worst of both.
- 3. State the direction of every measure.** The beginner could not tell whether "worst" was bad. Each factor knows whether high is good, because the factor library declares it. The interface does not say so.
- 4. Reconcile rarity against informativeness.** A finding can be unusual for a company and uninformative for its document type at the same time. The page should say which it is rather than printing both labels and leaving the reader to resolve it.
- 5. Keep broadening the read corpus.** Forty of a thousand and six companies are read, and the corpus behind the genre norms is still just over half technology. The drip now prefers the least covered sector, but it is bounded by a free tier and moves at roughly one company per run.

Limitations

Twenty-one paired decisions is a small sample and no significance is claimed for any figure here. The personas are language models rather than people. The seven companies are among the best covered in the database and are not representative of the other nine hundred and ninety-nine. Absolute values cannot be compared with the September run for the reasons given under Comparability. Most importantly, this trial measures influence on decisions and not the quality of those decisions, which remains unmeasured in both runs.